

TRACE NORM REGULARISED DEEP MULTI-TASK LEARNING

Yongxin Yang, Timothy M. Hospedales

Queen Mary, University of London

{yongxin.yang, t.hospedales}@qmul.ac.uk

ABSTRACT

We propose a framework for training multiple neural networks simultaneously. The parameters from all models are regularised by the tensor trace norm, so that each neural network is encouraged to reuse others' parameters if possible – this is the main motivation behind multi-task learning. In contrast to many deep multi-task learning models, we do not predefine a parameter sharing strategy by specifying which layers have tied parameters. Instead, our framework considers sharing for all shareable layers, and the sharing strategy is learned in a data-driven way.

1 INTRODUCTION AND RELATED WORK

Multi-task learning (MTL) (Caruana, 1997) aims to learn multiple tasks jointly, so that knowledge obtained from one task can be reused by others. We first briefly review some studies in this area.

Matrix-based Multi-Task Learning Matrix-based MTL is usually built on linear models, i.e., each task is parameterised by a D -dimensional weight vector w , and the model prediction is $\hat{y} = x \cdot w = x^T w$, where x is a D -dimensional feature vector representing an instance. The objective function for matrix-based MTL can be written as $\sum_{i=1}^T \sum_{j=1}^{N^{(i)}} \ell(y_j^{(i)}, x_j^{(i)} \cdot w^{(i)}) + \lambda \Omega(W)$. Here $\ell(y, \hat{y})$ is a loss function of the true label y and predicted label \hat{y} . T is the number of tasks, and for the i -th task there are $N^{(i)}$ training instances. Assuming the dimensionality of every task's feature is the same, the models – $w^{(i)}$ s – are of the same size. Then the collection of $w^{(i)}$ s forms a $D \times T$ matrix W of which the i -th column is the linear model for the i -th task. To achieve MTL we exploit a regulariser $\Omega(W)$ that couples the learning problems, typically by encouraging W to be a low-rank matrix. Some choices include the $\ell_{2,1}$ norm (Argyriou et al., 2008), and trace norm (Ji & Ye, 2009). An alternative approach (Kumar & Daumé III, 2012) is to explicitly formulate W as a low-rank matrix, i.e., $W = LS$ where L is a $D \times K$ matrix and S is a $K \times T$ matrix with $K < \min(D, T)$ as a hyper-parameter (matrix rank).

Tensor-based Multi-Task Learning In the classic MTL setting, each task is indexed by a single factor. But in many real-world problems, tasks are indexed by multiple factors. For example, to build a restaurant recommendation system, we want a regression model that predicts the scores for different aspects (food quality, environment) by different customers. Then the task is indexed by aspects \times customers. The collection of linear models for all tasks is then a 3-way tensor \mathcal{W} of size $D \times T_1 \times T_2$, where T_1 and T_2 is the number of aspects and customers respectively. Consequently $\Omega(\mathcal{W})$ has to be a tensor regulariser (Tomioka et al., 2010). For example, sum of the trace norms on all matriciations¹ (Romera-paredes et al., 2013), and scaled latent trace norm (Wimalawarne et al., 2014). An alternative solution is to concatenate the one-hot encodings of the two task factors and feed it as input into a two-branch neural network model (Yang & Hospedales, 2015).

Multi-Task Learning for Neural Networks With the success of deep learning, many studies have investigated deep multi-task learning. Zhang et al. (2014) use a convolutional neural network to find facial landmarks as well as recognise face attributes (e.g., emotions). Liu et al. (2015) propose a neural network for query classification and information retrieval (ranking for web search). A key commonality of these studies is that they use a predefined sharing strategy. A typical design is to use

¹Matriciation is also known as tensor unfolding or flattening.

the same parameters for the bottom layers of the deep neural network and task-specific parameters for the top layers. This kind of architecture can be traced back to 2000s (Bakker & Heskes, 2003). However, modern neural network architectures contain a large number of layers, which makes the decision of ‘*at which layer to split the neural network for different tasks?*’ extremely hard.

2 METHODOLOGY

Instead of predefining a parameter sharing strategy, we propose the following framework: For T tasks, each is modelled by a neural network of the same architecture. We collect the parameters in a layer-wise fashion, and put a tensor norm on every collection. We illustrate the idea by a simple example: assume that we have $T = 2$ tasks, and **each is modelled by a 4-layer convolution neural network (CNN). The CNN architecture is:** (1) convolutional layer (‘conv1’) of size $5 \times 5 \times 3 \times 32$, (2) ‘conv2’ of size $3 \times 3 \times 32 \times 64$, (3) fully-connected layer (‘fc1’) of size 256×256 , (4) fully-connected layer ‘fc2’⁽¹⁾ of size 256×10 for the first task and fully-connected layer (‘fc2’⁽²⁾) of size 256×20 for the second task. Since the two tasks have different numbers of outputs, the potentially shareable layers are ‘conv1’, ‘conv2’, and ‘fc1’, excluding the final layer of different dimensionality.

For single task learning, the parameters are ‘conv1’⁽¹⁾, ‘conv2’⁽¹⁾, ‘fc1’⁽¹⁾, and ‘fc2’⁽¹⁾ for the first task; ‘conv1’⁽²⁾, ‘conv2’⁽²⁾, ‘fc1’⁽²⁾, and ‘fc2’⁽²⁾ for the second task. We can see that there is not any parameter sharing between these two tasks. In one possible predefined deep MTL architecture, the parameters could be ‘conv1’, ‘conv2’, ‘fc1’⁽¹⁾, and ‘fc2’⁽¹⁾ for the first task; ‘conv1’, ‘conv2’, ‘fc1’⁽²⁾, and ‘fc2’⁽²⁾ for the second task, i.e., the first and second layer are fully shared in this case. For our proposed method, the parameter setting is the same as single task learning mode, but we put three tensor norms on the stacked {‘conv1’⁽¹⁾, ‘conv1’⁽²⁾} (a tensor of size $5 \times 5 \times 3 \times 32 \times 2$), the stacked {‘conv2’⁽¹⁾, ‘conv2’⁽²⁾} (a tensor of size $3 \times 3 \times 32 \times 64 \times 2$), and the stacked {‘fc1’⁽¹⁾, ‘fc1’⁽²⁾} (a tensor of size $256 \times 256 \times 2$) respectively.

Tensor Norm We choose to use the trace norm, the sum of a matrix’s singular values $\|X\|_* = \sum_{i=1} \sigma_i$. It has a nice property that it is the tightest convex relation of matrix rank (Recht et al., 2010). When directly restricting the rank of a matrix is challenging, trace norm serves as a good proxy. The extension of trace norm from matrix to tensor is not unique, just like tensor rank has multiple definitions. How to define tensor rank depends on how we assume the tensor is factorised, e.g., Tucker (Tucker, 1966) and Tensor-Train Oseledets (2011) decompositions. We propose three tensor trace norm designs here, which correspond to three variants of the proposed method.

For an N -way tensor \mathcal{W} of size $D_1 \times D_2 \times \dots \times D_N$. We define

$$\textbf{(Tensor Trace Norm) Last Axis Flattening} \quad \|\mathcal{W}\|_* = \gamma \|\mathcal{W}_{(N)}\|_* \quad (1)$$

$\mathcal{W}_{(i)} := \text{reshape}(\text{permute}(\mathcal{W}, [i, 1, \dots, i-1, i+1, \dots, N]), [D_i, \prod_{j \neq i} D_j])$ is the mode- i tensor flattening. This is the simplest definition. Given that in our framework, the last axis of tensor indexes the tasks, i.e., $D_N = T$, it is the most straightforward way to adapt the technique of matrix-based MTL – reshape the $D_1 \times D_2 \times \dots \times T$ tensor to $D_1 D_2 \dots \times T$ matrix.

To advance, we define two kinds of tensor trace norm that are closely connected with Tucker-rank (obtained by Tucker decomposition) and TT-rank (obtained by Tensor Train decomposition).

$$\textbf{(Tensor Trace Norm) Tucker} \quad \|\mathcal{W}\|_* = \sum_{i=1}^N \gamma_i \|\mathcal{W}_{(i)}\|_* \quad (2)$$

$$\textbf{(Tensor Trace Norm) TT} \quad \|\mathcal{W}\|_* = \sum_{i=1}^{N-1} \gamma_i \|\mathcal{W}_{[i]}\|_* \quad (3)$$

Here $\mathcal{W}_{[i]}$ is yet another way to unfold the tensor, which is obtained by $\mathcal{W}_{[i]} = \text{reshape}(\mathcal{W}, [D_1 D_2 \dots D_i, D_{i+1} D_{i+2} \dots D_N])$. It is interesting to note that unlike LAF, Tucker and TT also encourage within-task parameter sharing, e.g, sharing across filters in a neural network context.

Optimisation Using gradient-based methods for optimisation involving trace norm is *not* a common choice, as there are better solutions based on semi-definite programming or proximal gra-




Figure 1: Top-left: Testing accuracy. Top-mid: Training loss. Top-right: sharing strength by layer. Bottom: Norms when optimising LAF (left), Tucker (middle), TT (right).

dients since the trace norm is essentially non-differentiable. However, deep neural networks are usually trained by gradient descent, and we prefer to keep the standard training process. Therefore we use (sub-)gradient descent. The sub-gradient for trace norm can be derived as $\frac{\partial \|X\|_*}{\partial X} = X(X^T X)^{-\frac{1}{2}}$. A more numerical stable method instead of computing the inverse matrix square root is $X(X^T X)^{-\frac{1}{2}} = UV^T$ where U and V are obtained from SVD: $X = U\Sigma V^T$ (Watson, 1992).

3 EXPERIMENT

Our method is implemented in TensorFlow (Abadi et al., 2015), and released on Github². We experiment on the Omniglot dataset (Lake et al., 2015). Omniglot contains handwritten letters in 50 different alphabets (e.g., Cyrillic, Korean, Tengwar), each with its own number of unique characters (14 ~ 55). In total, there are 1623 unique characters, each with 20 instances. Each task is a multi-class character recognition problem for the corresponding alphabet. The images are monochrome of size 105×105 . We design a CNN with 3 convolutional and 2 FC layers. The first conv layer has 8 filters of size 5×5 ; the second conv layer has 12 filters of size 3×3 , and the third convolutional layer has 16 filters of size 3×3 . Each convolutional layer is followed by a 2×2 max-pooling. The first FC layer has 64 neurons, and the second FC layer has size corresponding to the number of unique classes in the alphabet. The activation function is *tanh*. We compare the three variants of the proposed framework – LAF (Eq. 1), Tucker (Eq. 2), and TT (Eq. 3) with single task learning (STL). For every layer, there are one (LAF) or more (Tucker and TT) γ that control the trade-off between the classification loss (cross-entropy) and the trace norm terms, for which we set all $\gamma = 0.01$.

The experiments are repeated 10 times, and every time 10% training data and 90% testing data are randomly selected. We plot the change of cross-entropy loss in training set and the values of norm terms with the neural networks’ parameters updating. As we can see in Fig 1, STL has the lowest training loss, but worst testing performance, suggesting over-fitting. Our methods alleviate the problem with multi-task regularisation. We roughly estimate the strength of parameter sharing by calculating $1 - \frac{\text{Norm of Optimised Param}}{\text{Norm of Initialised Param}}$, we can see the pattern that with bottom layers share more compared to the top ones. This reflects the common design intuition that the bottom layers are more data/task independent. Finally, it appears that the choice on LAF, Tucker, or TT may not be very sensitive as we observe that when optimising one, the loss of the other norms still reduces.

This technique provides a data-driven solution to the branching architecture design problem in deep multi-task learning. It is a flexible norm regulariser-based alternative to explicit factorisation-based approaches to the same problem (Yang & Hospedales, 2017).

²<https://github.com/wOOL/TNRDMTL>

REFERENCES

- Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dan Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. URL <http://tensorflow.org/>. Software available from tensorflow.org.
- Andreas Argyriou, Theodoros Evgeniou, and Massimiliano Pontil. Convex multi-task feature learning. *Machine Learning*, 2008.
- Bart Bakker and Tom Heskes. Task clustering and gating for Bayesian multitask learning. *Journal of Machine Learning Research (JMLR)*, 2003.
- Rich Caruana. Multitask learning. *Machine Learning*, 1997.
- Shuiwang Ji and Jieping Ye. An accelerated gradient method for trace norm minimization. In *International Conference on Machine Learning (ICML)*, 2009.
- Abhishek Kumar and Hal Daumé III. Learning task grouping and overlap in multi-task learning. In *International Conference on Machine Learning (ICML)*, 2012.
- Brenden M. Lake, Ruslan Salakhutdinov, and Joshua B. Tenenbaum. Human-level concept learning through probabilistic program induction. *Science*, 2015.
- Xiaodong Liu, Jianfeng Gao, Xiaodong He, Li Deng, Kevin Duh, and Ye-Yi Wang. Representation learning using multi-task deep neural networks for semantic classification and information retrieval. *NAACL*, 2015.
- I. V. Oseledets. Tensor-train decomposition. *SIAM Journal on Scientific Computing*, 2011.
- Benjamin Recht, Maryam Fazel, and Pablo A. Parrilo. Guaranteed minimum-rank solutions of linear matrix equations via nuclear norm minimization. *SIAM Rev.*, 2010.
- Bernardino Romera-paredes, Hane Aung, Nadia Bianchi-berthouze, and Massimiliano Pontil. Multilinear multitask learning. In *International Conference on Machine Learning (ICML)*, 2013.
- Ryota Tomioka, Kohei Hayashi, and Hisashi Kashima. On the extension of trace norm to tensors. In *NIPS Workshop on Tensors, Kernels, and Machine Learning*, 2010.
- L. R. Tucker. Some mathematical notes on three-mode factor analysis. *Psychometrika*, 1966.
- G.A. Watson. Characterization of the subdifferential of some matrix norms. *Linear Algebra and its Applications*, 170:33 – 45, 1992. ISSN 0024-3795. doi: [http://dx.doi.org/10.1016/0024-3795\(92\)90407-2](http://dx.doi.org/10.1016/0024-3795(92)90407-2). URL <http://www.sciencedirect.com/science/article/pii/0024379592904072>.
- Kishan Wimalawarne, Masashi Sugiyama, and Ryota Tomioka. Multitask learning meets tensor factorization: task imputation via convex optimization. In *Neural Information Processing Systems (NIPS)*, 2014.
- Yongxin Yang and Timothy M. Hospedales. A unified perspective on multi-domain and multi-task learning. In *International Conference on Learning Representations (ICLR)*, 2015.
- Yongxin Yang and Timothy M. Hospedales. Deep multi-task representation learning: A tensor factorisation approach. In *International Conference on Learning Representations (ICLR)*, 2017.
- Zhanpeng Zhang, Ping Luo, Chen Change Loy, and Xiaoou Tang. Facial landmark detection by deep multi-task learning. In *European Conference on Computer Vision (ECCV)*, 2014.